

Arrival Estimation in Emergency Departments

Context and Data



Emergency Room @ Maggiore Hospital

We will now consider a problem from the healthcare sector

We will use a dataset for the "Maggiore" hospital in Bologna



- In particular, we will focus on predicting arrivals
- ...To the Emergency Department (Pronto Soccorso)



A Look at the Dataset

We will start as usual by having a look at the dataset

```
In [59]: data = util.load_ed_data(data_file)
data
```

Out[59]:

	year	ID	Triage	TkCharge	Code	Outcome
0	2018	1	2018-01-01 00:17:33	2018-01-01 04:15:36	green	admitted
1	2018	2	2018-01-01 00:20:33	2018-01-01 03:14:19	green	admitted
2	2018	3	2018-01-01 00:47:59	2018-01-01 04:32:30	white	admitted
51238	2018	51239	2018-01-01 00:49:51	NaT	white	abandoned
51240	2018	51241	2018-01-01 01:00:40	NaT	green	abandoned
...
95665	2019	95666	2019-10-31 23:26:54	2019-10-31 23:41:13	yellow	admitted
95666	2019	95667	2019-10-31 23:46:43	2019-11-01 09:30:25	green	admitted
108622	2019	108623	2019-10-31 23:54:05	NaT	green	abandoned
95667	2019	95668	2019-10-31 23:55:32	2019-11-01 00:18:46	yellow	admitted
108623	2019	108624	2019-10-31 23:59:21	NaT	green	abandoned

108625 rows × 6 columns

A Look at the Dataset

Dataset fields and there are four relevant fields:


```
In [60]: data.iloc[:2]
```

```
Out[60]:
```

	year	ID	Triage	TkCharge	Code	Outcome
0	2018	1	2018-01-01 00:17:33	2018-01-01 04:15:36	green	admitted
1	2018	2	2018-01-01 00:20:33	2018-01-01 03:14:19	green	admitted

- Triage is the arrival time of each patient
- TKCharge is the time when a patient starts the first visit
- Code refers to the estimated priority (white < green < yellow < red)
- Outcome discriminates some special conditions (people quitting, fast tracks)


We'll sort the rows by increasing triage time:

```
 In [61]: data.sort_values(by='Triage', inplace=True)
```

Inter-Arrival Times

Let's check empirically the distribution of the inter-arrival times

```
In [62]: delta = data['Triage'].iloc[1:].round('2min') - data['Triage'].iloc[:-1].round('2min')
tmp = delta.value_counts().sort_index().values
tmp = pd.Series(index=np.arange(len(tmp)), data=tmp)
util.plot_bars(tmp, tick_gap=10, figsize=figsize)
```




There is a number of very low inter-arrival times

Waiting Time

Here is the distribution of the waiting times

```
In [63]: tmp = data[~data['TkCharge'].isnull()]
wait_time = tmp['TkCharge'].round('10min') - tmp['Triage'].round('10min')
tmp = wait_time.value_counts().sort_index().values
tmp = pd.Series(index=np.arange(len(tmp)), data=tmp)
util.plot_bars(tmp, tick_gap=10, figsize=figsize)
```



■ The distribution is **heavy-tailed** (large waiting times are quite relatively likely)

Arrival Estimation in Emergency Departments

Number of Arrivals



Binning

We will be interested in estimating the number of arrivals in a hour

First, we use a one-hot encoding for the priority codes

```
In [64]: codes = pd.get_dummies(data['Code'])
codes.set_index(data['Triage'], inplace=True)
codes.columns = codes.columns.to_list()
print(f'Number of examples: {len(codes)}')
codes.head()
```

Number of examples: 108625

Out[64]:

	green	red	white	yellow
Triage				
2018-01-01 00:17:33	True	False	False	False
2018-01-01 00:20:33	True	False	False	False
2018-01-01 00:47:59	False	False	True	False
2018-01-01 00:49:51	False	False	True	False
2018-01-01 01:00:40	True	False	False	False

Resampling

Then, we need to aggregate data with a specified frequency

```
In [65]: codes_b = codes.resample('H').sum()  
print(f'Number of examples: {len(codes_b)}')  
codes_b.head()
```

Number of examples: 16056

Out[65]:

	green	red	white	yellow
Triage				
2018-01-01 00:00:00	2	0	2	0
2018-01-01 01:00:00	7	1	1	1
2018-01-01 02:00:00	4	1	4	3
2018-01-01 03:00:00	7	0	1	1
2018-01-01 04:00:00	3	0	2	0

We count the arrivals in each hour, for each code



Computing Totals

Then we compute the total arrival counts

```
In [66]: cols = ['white', 'green', 'yellow', 'red']
codes_b['total'] = codes_b[cols].sum(axis=1)
codes_b
```

Out[66]:

	green	red	white	yellow	total
Triage					
2018-01-01 00:00:00	2	0	2	0	4
2018-01-01 01:00:00	7	1	1	1	10
2018-01-01 02:00:00	4	1	4	3	12
2018-01-01 03:00:00	7	0	1	1	9
2018-01-01 04:00:00	3	0	2	0	5
...
2019-10-31 19:00:00	3	1	0	4	8
2019-10-31 20:00:00	9	0	2	0	11
2019-10-31 21:00:00	3	0	0	2	5
2019-10-31 22:00:00	1	2	3	1	7
2019-10-31 23:00:00	5	0	0	2	7

16056 rows × 5 columns

Adding Time Information

Finally, we add time information (for later convenience)

```
In [67]: codes_bt = codes_b.copy()  
codes_bt['month'] = codes_bt.index.month  
codes_bt['weekday'] = codes_bt.index.weekday  
codes_bt['hour'] = codes_bt.index.hour  
codes_bt
```

Out[67]:

	green	red	white	yellow	total	month	weekday	hour
Triage								
2018-01-01 00:00:00	2	0	2	0	4	1	0	0
2018-01-01 01:00:00	7	1	1	1	10	1	0	1
2018-01-01 02:00:00	4	1	4	3	12	1	0	2
2018-01-01 03:00:00	7	0	1	1	9	1	0	3
2018-01-01 04:00:00	3	0	2	0	5	1	0	4
...
2019-10-31 19:00:00	3	1	0	4	8	10	3	19
2019-10-31 20:00:00	9	0	2	0	11	10	3	20
2019-10-31 21:00:00	3	0	0	2	5	10	3	21
2019-10-31 22:00:00	1	2	3	1	7	10	3	22




Counts over Time

Our resampled series can be plotted easily over time

Let's see the total counts as an example:

```
In [68]: util.plot_series(codes_b['total'], figsize=figsize)
```




Variability

With our binned series, we can assess the count variability

Let's check it over different months:

```
In [69]: codes_bt[['month', 'total']].boxplot(by='month', figsize=figsize);
```




■ The variability does not change much over different months

Variability

Here is the standard deviation over weekdays

```
In [70]: codes_bt[['weekday', 'total']].boxplot(by='weekday', figsize=figsize);
```




- There is a trend, but rather weak



Variability

...And finally over hours

```
In [71]: codes_bt[['hour', 'total']].boxplot(by='hour', figsize=figsize);
```



- Variance and mean seem to be quite correlated



Arrival Estimation in Emergency Departments

Estimating Stochastic Quantities



Arrival Prediction

We can now frame our arrival prediction problem

We have some **input information**:

- Hour, day of the week, and month
- ...Plus possibly the observed arrivals in previous hours

We want to **predict the number of arrivals** in the next interval

Have we encountered similar tasks in other use cases?



Arrival Prediction

We can now frame our arrival prediction problem

We have some **input information**:

- Hour, day of the week, and month
- ...Plus possibly the observed arrivals in previous hours

We want to **predict the number of arrivals** in the next interval

Have we encountered similar tasks in other use cases?

On the face of it, this is a **regression problem**

But there is a catch!




Prediction and Randomness

The number of arrivals is not subject to a lot of uncertainty!

Let's check its values against the most informative input, i.e. the hour of the day

```
In [72]: tmp = codes_b[codes_b.index.hour == 6]['total']  
tmpv = tmp.value_counts(sort=False, normalize=True).sort_index()  
util.plot_bars(tmpv, figsize=figsize)
```



■ There isn't a single, very likely value: the number of arrival is **stochastic**!

Identifying the Distribution

Instead of predicting a value, we can predict **the probability of every possible value**

Formally, our goal is estimating a **conditional distribution**

$$P(Y \mid X)$$

- X is the observable input information we choose to employ
- Y is the number of arrivals in the next hour



Identifying the Distribution

Instead of predicting a value, we can predict **the probability of every possible value**

Formally, our goal is estimating a **conditional distribution**

$$P(Y \mid X)$$

- X is the observable input information we choose to employ
- Y is the number of arrivals in the next hour

We can think of training a parameterized model on this purpose

$$\hat{f}(x; \theta) \simeq P(Y \mid X)$$

- We will see one viable approach to achieve that
- ...Provided that we know the type of distribution we want to predict



Poisson Distribution

Many arrival process are well described by Poisson distributions

The Poisson distribution is defined by a single parameter λ

λ is the rate of occurrence of the events

- The distribution has a **discrete support**
- The Probability Mass Function is:

$$p(k, \lambda) = \frac{\lambda^k e^{-\lambda}}{k!}$$

- Both the **mean** and the **standard deviation** have the same value (i.e. λ)

The distribution is a good choice provided that the events we are counting are:


- Independent
- Happening with a constant rate



Fitted Poisson Distribution

Let's try to fit a Poisson distribution over our target

```
In [74]: mu = tmp.mean()  
dist = stats.poisson(mu)  
x = np.arange(tmp.min(), tmp.max()+1)  
util.plot_bars(tmpv, figsize=figsize, series=pd.Series(index=x, data=dist.pmf(x)))
```




It's a very good match!

Fitted Poisson Distribution

Let's try for 8AM (closer to the peak)


```
In [75]: tmp = codes_b[codes_b.index.hour == 8]['total']  
tmpv = tmp.value_counts(sort=False, normalize=True).sort_index()  
mu = tmp.mean()  
dist = stats.poisson(mu)  
x = np.arange(tmp.min(), tmp.max()+1)  
util.plot_bars(tmpv, figsize=figsize, series=pd.Series(index=x, data=dist.pmf(x)))
```



Fitted Poisson Distribution

...And finally for the peak itself (11am)

```
In [76]: tmp = codes_b[codes_b.index.hour == 11]['total']  
tmpv = tmp.value_counts(sort=False, normalize=True).sort_index()  
mu = tmp.mean()  
dist = stats.poisson(mu)  
x = np.arange(tmp.min(), tmp.max()+1)  
util.plot_bars(tmpv, figsize=figsize, series=pd.Series(index=x, data=dist.pmf(x)))
```





Learning and Estimator

How can we build an estimator for our problem?



Learning and Estimator

How can we build an estimator for our problem?

We could build a table

For example, we could compute average arrivals for every hour of the day

- These correspond to λ for that hour, so we target the correct distribution
- ...But the approach has trouble scaling to multiple features



Learning and Estimator

How can we build an estimator for our problem?

We could build a table

For example, we could compute average arrivals for every hour of the day

- These correspond to λ for that hour, so we target the correct distribution
- ...But the approach has trouble scaling to multiple features

We could train a regressor as usual

For example a Linear Regressor or a Neural Network, with the classical MSE loss

- ✎ If we do this, it's easy to include multiple input features

Neuro-Probabilistic Models

In practice there is an alternative

Let's start by build a **probabilistic model** of our phenomenon:

$$y \sim \text{Pois}(\lambda(x))$$

- The number arrivals in a 1-hour bin (i.e. y)
- ...Is **drawn from a Poisson distribution** (parameterized with a rate)
- ...But **the rate is a function** of known input, i.e. $\lambda(x)$



Neuro-Probabilistic Models

In practice there is an alternative

Let's start by build a **probabilistic model** of our phenomenon:

$$y \sim \text{Pois}(\lambda(x))$$

- The number arrivals in a 1-hour bin (i.e. y)
- ...Is **drawn from a Poisson distribution** (parameterized with a rate)
- ...But **the rate is a function** of known input, i.e. $\lambda(x)$

Then we can approximate lambda using an estimator, leading to:

$$y \sim \text{Pois}(\lambda(x, \theta))$$

- $\lambda(x, \theta)$ can be any model, with parameter vector λ

This is a **hybrid** approach, combining statistics and ML

Neuro-Probabilistic Models

How do we train this kind of model?

Just as usual, i.e. for (empirical) maximum log likelihood:

$$\operatorname{argmin}_{\theta} - \sum_{i=1}^m \log f(\hat{y}_i, \lambda(\hat{x}_i, \theta))$$

- Where $f(\hat{y}_i, \lambda)$ is the probability of value \hat{y}_i according to the distribution
- ...And $\lambda(\hat{x}_i, \theta)$ is the estimate rate for the input \hat{x}_i

In detail, in our case we have:

$$\operatorname{argmin}_{\theta} - \sum_{i=1}^m \log \frac{\lambda(\hat{x}_i, \theta)^{\hat{y}_i} e^{-\lambda(\hat{x}_i, \theta)}}{\hat{y}_i!}$$

 ...Which is differentiable and can be solved via gradient descent!

Building a Neuro-Probabilistic Model

We can build this class of models by using custom loss functions

...But it's easier to use a library such as TensorFlow Probability

- TFP provides a layer that abstracts a generic probability distribution:

```
tfp.layers.DistributionLambda(distribution_function, ...)
```

- And function (classes) to model many statistical distributions, e.g.:

```
tfp.distributions.Poisson(log_rate=None, ...)
```

About the **DistributionLambda** layer

- Its input is a symbolic tensor (like for any other layer)
- Its output is tensor of probability distribution **objects**
- ...Rather than a tensor of numbers



Building a Neuro-Probabilistic Model

The `util` module contains code to build our neuro-probabilistic model

```
def build_nn_poisson_model(input_shape, hidden, rate_guess=1):
    model_in = keras.Input(shape=input_shape, dtype='float32')
    x = model_in
    for h in hidden:
        x = layers.Dense(h, activation='relu')(x)
    log_rate = layers.Dense(1, activation='linear')(x)
    lf = lambda t: tfp.distributions.Poisson(rate=rate_guess * tf.math.exp(t))
    model_out = tfp.layers.DistributionLambda(lf)(log_rate)
    model = keras.Model(model_in, model_out)
    return model
```

- An MLP architecture computes the `log_rate` tensor (corresponding to $\log \lambda(x)$)
- Using a log, we make sure the rate is **strictly positive**
- A `DistributionLambda` yield the output (a distribution object)



Building a Neuro-Probabilistic Model

The `util` module contains code to build our neuro-probabilistic model

```
def build_nn_poisson_model(input_shape, hidden, rate_guess=1):
    model_in = keras.Input(shape=input_shape, dtype='float32')
    x = model_in
    for h in hidden:
        x = layers.Dense(h, activation='relu')(x)
    log_rate = layers.Dense(1, activation='linear')(x)
    lf = lambda t: tfp.distributions.Poisson(rate=rate_guess * tf.math.exp(t))
    model_out = tfp.layers.DistributionLambda(lf)(log_rate)
    model = keras.Model(model_in, model_out)
    return model
```

- The `DistributionLambda` layer is parameterized with a function
- The function (`lf` in this case) constructs the distribution object
- ...Based on its input tensor (called `t` in the code)



Building a Neuro-Probabilistic Model

We need to be careful about **initial parameter estimates**

```
def build_nn_poisson_model(input_shape, hidden, rate_guess=1):  
    ...  
    lf = lambda t: tfp.distributions.Poisson(rate=rate_guess * tf.math.exp(t))  
    ...
```

- Assuming standardized/normalized input, under default weight initialization
- ...The `log_rate` tensor will be initially close to 0
- Meaning out rate λ would be initially close to $e^0 = 1$

We need to make sure that this guess is **meaningful for our target**

- In the code, this is achieved by scaling the rate
- ...With a guess that must be passed at model construction time



Training a Neuro-Probabilistic Model

Training the model requires to specify the loss function

...Which in our case is the **negative log-likelihood**

- So, it turns out we do need a custom loss functions
- ...But with TFP this is easy to compute

In particular, as loss function we **always** use:

```
negloglikelihood = lambda y_true, dist: -dist.log_prob(y_true)
```

- The first parameter is the observed value (e.g. actual number of arrivals)
- The second is the distribution computed by the `DistributionLambda` layer
- ...Which provides the method `log_prob`



Data Preparation

Let's see the approach in practice

We will start by preparing our data:

- As input we will use the field weekday in natural form
- ...And the field hour using a one-hot encoding

Let's perform the encoding:

```
In [77]: np_data = pd.get_dummies(codes_bt, columns=['hour'], dtype='int32')
np_data.iloc[:2]
```

Out[77]:

	green	red	white	yellow	total	month	weekday	hour_0	hour_1	hour_2	...	hour_14	hour_15	hour_16	hour_17
Triage															
2018-01-01 00:00:00	2	0	2	0	4	1	0	1	0	0	...	0	0	0	0
2018-01-01 01:00:00	7	1	1	1	10	1	0	0	1	0	...	0	0	0	0

Data Preparation

Now we can separate the training and test data

```
In [78]: sep = '2019-01-01'
np_tr = np_data[np_data.index < sep]
np_ts = np_data[np_data.index >= sep]
```

...And then the input and output

```
In [79]: in_cols = [c for c in np_data.columns if c.startswith('hour')] + ['weekday']
out_col = 'total'

np_tr_in = np_tr[in_cols].copy()
np_tr_in['weekday'] = np_tr_in['weekday'] / 6
np_tr_out = np_tr[out_col].astype('float64')

np_ts_in = np_ts[in_cols].copy()
np_ts_in['weekday'] = np_ts_in['weekday'] / 6
np_ts_out = np_ts[out_col].astype('float64')
```



Data Preparation

The input data need to be standardized/normalized as usual

In our case, we do this only for weekday (the hours are already $\in \{0, 1\}$)

```
np_tr_in['weekday'] = np_tr_in['weekday'] / 6
```

The output does not require standarization

...But we need to represent it using floating point numbers

```
np_tr_out = np_tr[out_col].astype('float64')
```

- This is an implementation requirement for TensorFlow




Building the Model

We can now build the Neuro-Probabilistic model


```
In [80]: nnp = util.build_nn_poisson_model(input_shape=len(in_cols), hidden=[], rate_guess=np_t
util.plot_nn_model(nnp)
```

Out[80]:



As a rate guess we use the average over the training set

- This is easy to compute
- ...And will provide a better starting point for gradient descent




Training the Model

We can train the model (mostly) as usual

...Except that we need to use the mentioned custom loss function

```
In [81]: negloglikelihood = lambda y_true, dist: -dist.log_prob(y_true)
nnp = util.build_nn_poisson_model(input_shape=len(in_cols), hidden=[], rate_guess=np_t
history = util.train_nn_model(nnp, np_tr_in, np_tr_out, loss=negloglikelihood, validat
util.plot_training_history(history, figsize=figsize)
```



Final loss: 2.2627 (training)

Predictions

When we call the `predict` method on the model we obtain **samples**

This means that the result of `predict` is **stochastic**

```
In [82]: print(str(nnp.predict(np_tr_in, verbose=0)[:3]).replace('\n', ' '))  
         print(str(nnp.predict(np_tr_in, verbose=0)[:3]).replace('\n', ' '))
```

```
[[1.] [7.] [4.]]  
[[3.] [4.] [5.]]
```

We can obtain the distribution object by simply **calling the model**

```
In [83]: nnp(np_tr_in.values)
```

```
Out[83]: <tfp.distributions._TensorCoercible 'tensor_coercible' batch_shape=[8760, 1] event_shape=[] dtype=float32>
```



Then we can call **methods over the distribution objects**


Evaluation

Using the predict means, let's check the quality of our results

```
In [84]: tr_pred = nnp(np_tr_in.values).mean().numpy().ravel()  
util.plot_pred_scatter(np_tr_out, tr_pred, figsize=figsize)
```

R2: 0.60

MAE: 1.93



This is a **stochastic** process, making this R^2 value **very good**


Evaluation

Let's repeat the exercise on the test set

```
In [85]: ts_pred = nnp(np_ts_in.values).mean().numpy().ravel()  
util.plot_pred_scatter(np_ts_out, ts_pred, figsize=figsize)
```

R2: 0.60

MAE: 1.94



No overfitting, which is again very good

Confidence Intervals

Since our output is a distribution, we have access to **all sort of statistics**

Here we will simply show the mean and stdev over one week of data:

```
In [89]: ts_pred_std = nnp(np_ts_in.values).stddev().numpy().ravel()  
util.plot_series(pd.Series(index=np_ts_in.index[:24*7], data=ts_pred[:24*7]), std=pd.S  
plt.scatter(np_ts_in.index[:24*7], np_ts_out[:24*7], marker='x');
```

